

Learning Structure in Time with a Plastic Recurrent Neural Network

Fabian Schubert

February 21, 2018

1 Introduction

We implemented a neural network consisting of binary neurons, modeled in discrete time steps, which follows the ideas presented in [1]. The architecture of our network is depicted in Fig. 1 and shall be described in further detail.

The recurrent network consists of a population with $N_e = 300$ excitatory units (denoted as x_e) and a population with $N_i = 60$ (denoted as x_i) inhibitory units. Furthermore, a population of $N_{ext} = 9$ excitatory units (I_j) is interpreted as external input, where the input coming from each external unit is to be interpreted as encoding a particular feature of a sensory stream, e.g. the recognition of a particular letter or symbol.

2 Methods

2.1 Network Details

Synaptic connectivities - represented by arrows in the illustration - were initially generated from a uniform distribution and the following properties, listed in Table 1.

2.2 Neuron Model

The state of the neurons is updated in discrete time steps by the following equations:

Table 1: Network parameters	
Connection Fraction W_{ee}	0.05
Connection Fraction W_{ei}	0.1
Connection Fraction W_{ie}	0.2
Connection Fraction W_{ii}	0.2
Connection Fraction $W_{e,ext}$	0.1
$\langle W_{e,ext} \rangle$	0.2
Total postsynaptic E→E input	1.0
Total postsynaptic I→E input	-1.0
Total postsynaptic E→I input	1.0
Total postsynaptic I→I input	-1.0

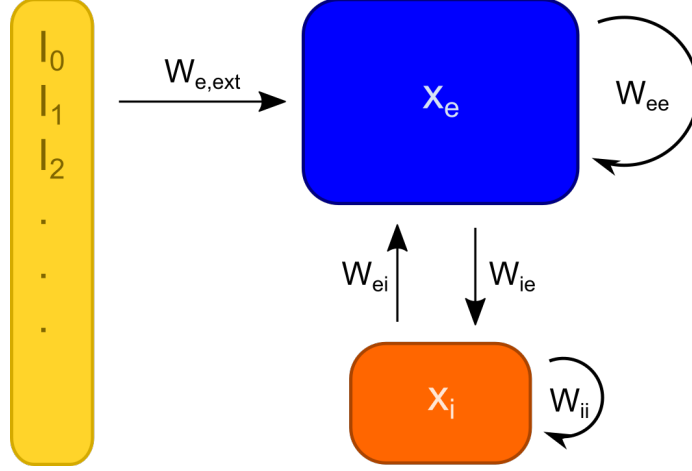


Figure 1: Architecture of the RNN

Table 2: Parameters of the neuron model

$\mu_{noise,e}$	0.0
$\mu_{noise,i}$	0.0
$\sigma_{noise,e}$	0.1
$\sigma_{noise,i}$	0.1
$\mu_{target,e}$	0.1
$\mu_{target,i}$	0.1
$\sigma_{target,e}$	0.0
$\sigma_{target,i}$	0.0
μ_{IP}	0.002

$$x_{e,n}(t+1) = \theta \left(\sum_{j=0}^{N_e-1} W_{ee,nj} x_{e,j}(t) + \sum_{k=0}^{N_i-1} W_{ei,nk} x_{i,k}(t) + \sum_{l=0}^{N_{ext}-1} W_{e,ext,nl} I_l(t) - T_{e,n}(t) + \xi_{e,n}(t) \right) \quad (1)$$

$$x_{i,n}(t+1) = \theta \left(\sum_{j=0}^{N_e-1} W_{ie,nj} x_{e,j}(t) + \sum_{k=0}^{N_i-1} W_{ii,nk} x_{i,k}(t) - T_{i,n}(t) + \xi_{i,n}(t) \right) \quad (2)$$

where $\theta(\cdot)$ is the theta function and T_e and T_i represent additional threshold values. $\xi_{e/i}$ are random noise terms sampled from a Gaussian distribution at each time step with parameters $\mu_{noise,e/i}$ and $\sigma_{noise,e/i}$.

To stabilize network activity, each neuron's threshold is updated each time step such that the neuron's average activity approach a given target value:

$$T_{e/i,n}(t+1) = T_{e/i,n}(t) + \mu_{IP} (x_{e/i,n}(t) - r_{target,e/i,n}) \quad (3)$$

where μ_{IP} is the learning rate of this ‘‘intrinsic plasticity’’. Target rates $r_{target,e/i}$ were drawn randomly from a Gaussian distribution with parameters $\mu_{target,e/i}$ and $\sigma_{target,e/i}$ for each neuron and kept fixed throughout the simulation.

Parameters of the dynamics described in this section are given in Table 2.




Figure 2: Transition Matrix between subsequently active input states

2.3 Rules of the Input Sequence

In each time step, only a single unit is in its active state $I_j(t) = 1$. The sequence of active input nodes was generated by a markov chain with transition probabilities shown in Fig. 2.

Due to the structure of the transition matrix, the sequence is partially predictable in the sense that an element of $\{0, 1, 2\}$ will always be followed by an element of $\{3, 4, 5\}$ etc.

2.4 Plasticity Rules

Recurrent excitatory connection were subject to two plasticity mechanisms: A simple pre-post Hebbian learning rule and a postsynaptic multiplicative normalization preventing connectivity runaway.

$$\Delta W_{ee,ij}(t) = \mu_{hebb} (x_{e,j}(t-1)x_{e,i}(t) - x_{e,i}(t-1)x_{e,j}(t)) \quad (4)$$

$$W_{ee,ij}(t) = w_{total,ee} \frac{W_{ee,ij}(t-1) + \Delta W_{ee,ij}(t)}{\sum_{k=0}^{N_e-1} W_{ee,ik}(t-1) + \Delta W_{ee,ik}(t)} \quad (5)$$

We did not include pruning or creation of synapses, but set a very small lower bound for existing excitatory connections.

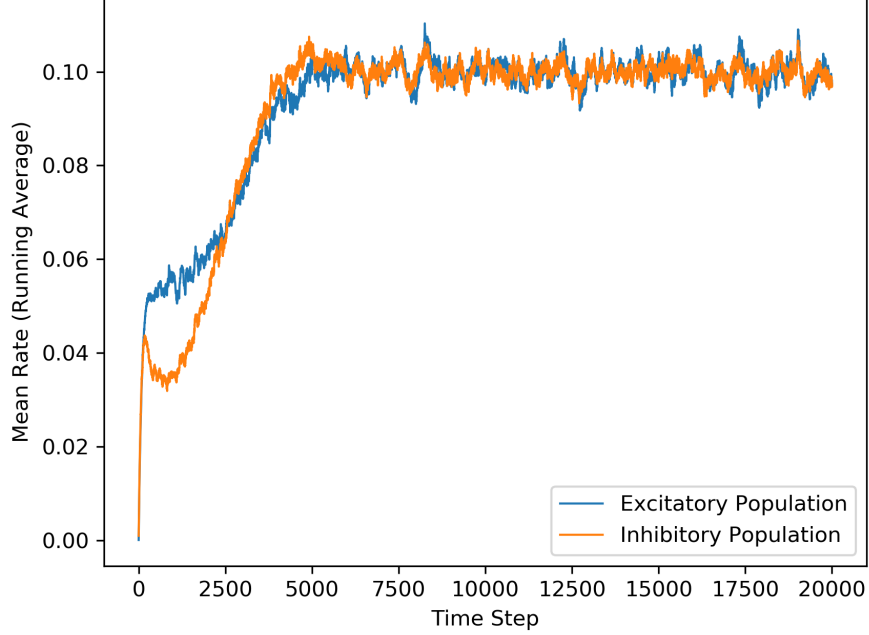


Figure 3: Running average of excitatory and inhibitory population activity.

3 Results

Network activity settled at a constant mean rate and a corresponding mean threshold, see Fig. 3 and Fig. 4.

Furthermore, Fig. 5 and Fig. 6 suggest that the appearance of active states follows poissonian statistics. However, the distribution of excitatory inter“spike”-intervals shows a clear preference for multiples of 3, which was not present in the absence of external input and is reflected in the 3-fold periodicity of the input sequence.

Generally speaking, the implemented plasticity rules often gave rise to time courses of synaptic weights similar to the one shown in Fig. 7: the emergence of one or more comparably strong weights alongside a majority of weak connections.

3.1 Cluster Analysis of Excitatory Activity

Following the conceptual idea presented by Elman [2], we performed a hierarchical cluster analysis of the binary activity vectors of the excitatory population. For this, we used the activity data of x_e from the last 3000 steps. We then performed a hierarchical cluster analysis with Ward’s method. The resulting dendrogram is depicted in Fig. 8. A 3-fold structure is visible in the uppermost branching layer.




Figure 4: Population mean of excitatory and inhibitory thresholds.

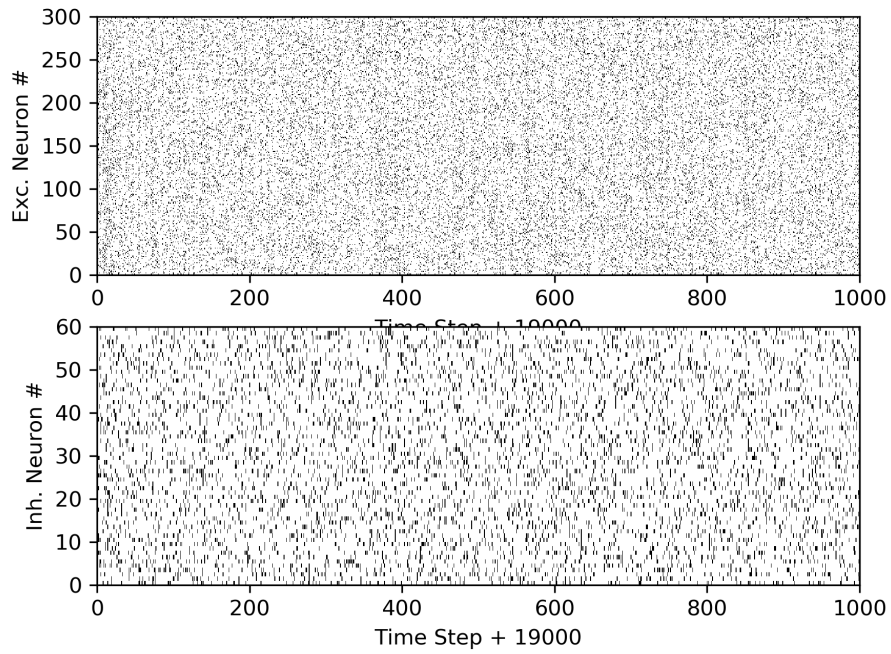


Figure 5: Raster plot of network activity.

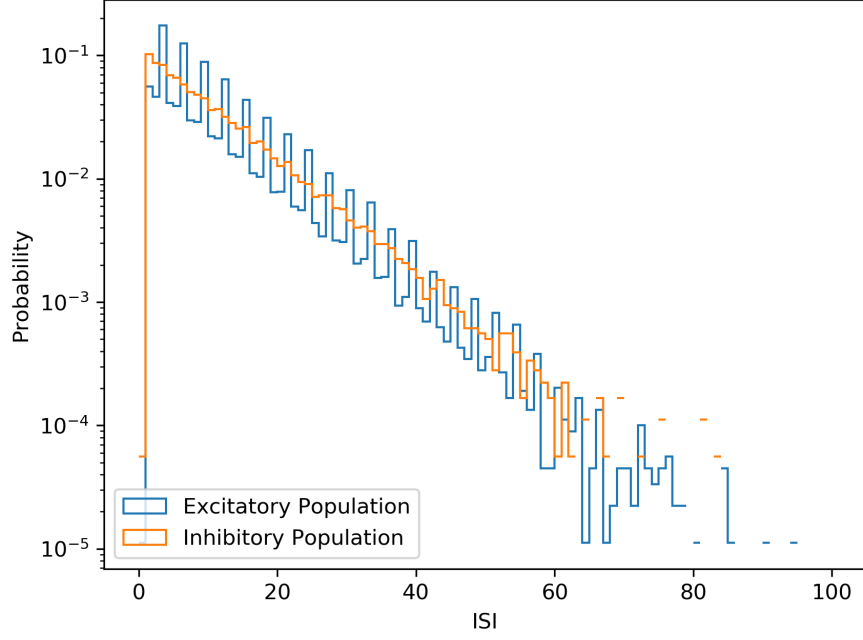


Figure 6: Distribution of interspike intervals.




Figure 7: Sample time course of E-iE weights.

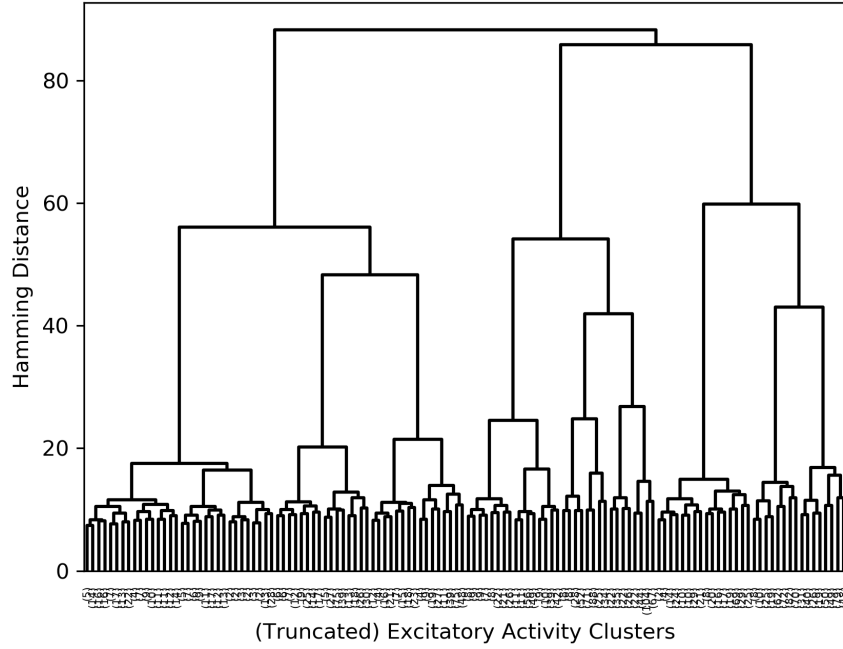


Figure 8: Dendrogram of a cluster analysis of excitatory activity patterns.

References

- [1] R. Duarte, P. Series, and A. Morrison. Self-Organized Artificial Grammar Learning in Spiking Neural Networks. In *36th Annual Conference of the Cognitive Science Society*, 07 2014.
- [2] J. L. Elman. Finding Structure in Time. *Cognitive Science*, 14(2):179–211, 1990.